

[CM2] Our Own Network

```
In [ ]: class Simple_CNN(nn.Module):
    def __init__(self, out_1=16, out_2=32, out_3 = 64):
        super(Simple_CNN, self).__init__()
        self.cnn1 = nn.Conv2d(in_channels=1, out_channels=out_1, kernel_size=3, stride=1)
        self.relu1 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.cnn2 = nn.Conv2d(in_channels=out_1, out_channels=out_2, kernel_size=3, stride=1)
        self.relu2 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.cnn3 = nn.Conv2d(in_channels=out_2, out_channels=out_3, kernel_size=3, stride=1)
        self.maxpool3 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.relu3 = nn.ReLU()
        self.fc1 = nn.Linear(out_3 * 8 * 8, 500)
        self.fc2 = nn.Linear(500, 50)
        self.fc3 = nn.Linear(50, 5)
        self.softmax = nn.Softmax()
        self.sig = nn.Sigmoid()

    def forward(self,x):
        out = self.cnn1(x)
        out = self.relu1(out)
        out = self.maxpool1(out)
        out = self.cnn2(out)
        out = self.relu2(out)
        out = self.maxpool2(out)
        out = out.view(out.size(0), -1)
        out = self.fc1(out)
        out = self.relu1(out)
        out = self.fc2(out)
        out = self.relu1(out)
        out = self.fc3(out)
        out = self.softmax(out)
        return out

    def forward2(self,x):
        out = self.cnn1(x)
        out = self.relu1(out)
        out = self.maxpool1(out)
        out = self.cnn2(out)
        out = self.relu2(out)
        out = self.maxpool2(out)
        out = out.view(out.size(0), -1)
        out = self.fc1(out)
        out = self.relu1(out)
        out = self.fc2(out)
        out = self.relu1(out)
        out_enc = out.clone().detach()
        return out_enc
```

```

def train(self, train_loader,opt, loss_fn, BATCH_SIZE, eval = True):
    loss_tot = 0
    acc = 0
    cnt = 0
    for x,y in zip(train_loader[0], train_loader[1]):
        cnt+=1
        x = x.float()
        x = x.cuda()
        y = y.cuda()
        opt.zero_grad()
        z = self.forward(x)
        y = torch.nn.functional.one_hot(y, num_classes=5)
        loss = loss_fn(z.double(),y.double())
        loss_tot+= loss.item()
        loss.backward()
        opt.step()
    loss_tr = loss_tot/(len(train_loader[1]))
    if eval:
        correct = 0
        for x,y in zip(train_loader[0],train_loader[1]):
            x = x.float()
            x = x.cuda()
            y = y.cuda()
            z = self.forward(x)
            _,yhat = torch.max(z.data,1)
            correct += (yhat== y).sum().item()
        acc = correct/(len(train_loader[1])*BATCH_SIZE)
    return loss_tr,acc

def eval(self,loader,BATCH_SIZE,loss_fn=nn.CrossEntropyLoss()):
    correct = 0
    acc = 0
    loss_tot = 0
    for x,y in zip(loader[0],loader[1]):
        x = x.float()
        x = x.cuda()
        y = y.cuda()
        z = self.forward(x)
        loss = loss_fn(z.double(),y.long())
        loss_tot += loss.item()
        _ , yhat = torch.max(z.data,1)
        correct += (yhat == y).sum().item()
    acc = correct/(len(loader[1])*BATCH_SIZE)
    loss_val = loss_tot/(len(loader[1])*BATCH_SIZE)
    return loss_val, acc

def pred(self,x):
    x = x.float()
    x = x.cuda()
    z = self.forward(x)
    _,yhat = torch.max(z.data,1)
    return yhat.item()

```

```

In [ ]: model2 = Simple_CNN()
        model2 = model2.cuda()
        summary(model2, (1,32,32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 32, 32]	160
ReLU-2	[-1, 16, 32, 32]	0
MaxPool2d-3	[-1, 16, 16, 16]	0
Conv2d-4	[-1, 32, 16, 16]	4,640
ReLU-5	[-1, 32, 16, 16]	0
MaxPool2d-6	[-1, 32, 8, 8]	0
Linear-7	[-1, 500]	1,024,500
ReLU-8	[-1, 500]	0
Linear-9	[-1, 50]	25,050
ReLU-10	[-1, 50]	0
Linear-11	[-1, 5]	255
Softmax-12	[-1, 5]	0
Total params: 1,054,605		
Trainable params: 1,054,605		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.43		
Params size (MB): 4.02		
Estimated Total Size (MB): 4.46		

Architecture:

The network architecture summary is shown above. We have 2 convolution layers, each with stride 1 and padding 1. First layer has 16 filters while the second one has 32 filters. The choice of pooling is max pool for all conv layers. The choice of activation function is ReLU. We flatten the output of convolution and employ FCC with 500 followed by a layer containing 50 neurons followed by the output layer. Here too, the choice of activation function is ReLU followed by Softmax at the output layer since our problem is of classification nature.

Optimizer: We have used "adam" optimizer while compiling the model because Adam combines the best properties of the AdaGrad and RMSProp algorithms to provide an optimization algorithm that can handle sparse gradients on noisy problems. Adam is relatively easy to configure where the default configuration parameters do well on most problems.

Loss Function: We used 'CrossEntropyLoss' since we one hot encoded the output.

```

In [ ]: opt = Adam(model2.parameters(), lr = 0.001)
        crit = nn.CrossEntropyLoss()
        model2 = model2.cuda()
        tr_loss2 = []
        tr_acc2 = []
        val_acc2 = []
        val_loss2 = []
        training_time_tot = 0
        for epoch in range(75):
            start_time = time.time()
            l, a = model2.train(train_loader=loader, opt = opt, loss_fn = crit, BATC
            end_time = time.time()
            training_time_tot += end_time - start_time
            print("Epoch ", epoch+1, "--> ", a)
            tr_loss2.append(l)
            tr_acc2.append(a)
            l, a = model2.eval(loader=val_loader, BATCH_SIZE=100)
            val_acc2.append(a)
            val_loss2.append(l)

```

```

Epoch 1 --> 0.7931041666666667
Epoch 2 --> 0.8113541666666667
Epoch 3 --> 0.8221875
Epoch 4 --> 0.8348958333333333
Epoch 5 --> 0.8303958333333333
Epoch 6 --> 0.8438541666666667
Epoch 7 --> 0.8510208333333333
Epoch 8 --> 0.8503541666666666
Epoch 9 --> 0.8568958333333333
Epoch 10 --> 0.8565625
Epoch 11 --> 0.86025
Epoch 12 --> 0.8541458333333334
Epoch 13 --> 0.8589583333333334
Epoch 14 --> 0.86375
Epoch 15 --> 0.8653541666666666
Epoch 16 --> 0.8458333333333333
Epoch 17 --> 0.8689166666666667
Epoch 18 --> 0.8700625
Epoch 19 --> 0.8732708333333333
Epoch 20 --> 0.8670416666666667
Epoch 21 --> 0.87075
Epoch 22 --> 0.8787708333333333
Epoch 23 --> 0.8764791666666667
Epoch 24 --> 0.875875
Epoch 25 --> 0.8773958333333334
Epoch 26 --> 0.8699791666666666
Epoch 27 --> 0.8807291666666667
Epoch 28 --> 0.8785416666666667
Epoch 29 --> 0.8787291666666667
Epoch 30 --> 0.8834375
Epoch 31 --> 0.87875
Epoch 32 --> 0.88175
Epoch 33 --> 0.8771875
Epoch 34 --> 0.8851041666666667
Epoch 35 --> 0.8830625
Epoch 36 --> 0.8865416666666667

```

```

Epoch 37 --> 0.8841666666666667
Epoch 38 --> 0.8883541666666667
Epoch 39 --> 0.8817916666666666
Epoch 40 --> 0.8891041666666667
Epoch 41 --> 0.8907083333333333
Epoch 42 --> 0.8871041666666667
Epoch 43 --> 0.8895416666666667
Epoch 44 --> 0.8890833333333333
Epoch 45 --> 0.88825
Epoch 46 --> 0.8867708333333333
Epoch 47 --> 0.8877916666666666
Epoch 48 --> 0.8911041666666667
Epoch 49 --> 0.8931041666666667
Epoch 50 --> 0.8928333333333334
Epoch 51 --> 0.8797291666666667
Epoch 52 --> 0.8860833333333333
Epoch 53 --> 0.8936875
Epoch 54 --> 0.8917083333333333
Epoch 55 --> 0.8851875
Epoch 56 --> 0.8945
Epoch 57 --> 0.8961458333333333
Epoch 58 --> 0.8908958333333333
Epoch 59 --> 0.8943333333333333
Epoch 60 --> 0.8954791666666667
Epoch 61 --> 0.8943958333333333
Epoch 62 --> 0.8958125
Epoch 63 --> 0.8902291666666666
Epoch 64 --> 0.8978333333333334
Epoch 65 --> 0.896375
Epoch 66 --> 0.8942708333333333
Epoch 67 --> 0.8981458333333333
Epoch 68 --> 0.8988125
Epoch 69 --> 0.8974375
Epoch 70 --> 0.8980625
Epoch 71 --> 0.8959791666666667
Epoch 72 --> 0.8953958333333333
Epoch 73 --> 0.898
Epoch 74 --> 0.8979583333333333
Epoch 75 --> 0.9007083333333333

```

```


In [ ]: plt.plot(tr_acc2, label = 'Training Accuracy')
plt.plot(val_acc2, label = 'Validation Accuracy')
plt.xlabel("Epoch")
plt.ylabel("Training Accuracy")
plt.title("Training Accuracy by Epoch plot - Model 2")
plt.legend()

```

```


Out[ ]: <matplotlib.legend.Legend at 0x7ff48b4c86d0>

```



```
In [ ]: plt.plot(np.array(tr_loss2)/100, label = 'Training Loss')
plt.plot(val_loss2, label = 'Validation Loss')
plt.xlabel("Epoch")
plt.ylabel("Training Loss")
plt.title("Training Loss by Epoch plot - Model 2")
plt.legend()
```

```
Out[ ]: <matplotlib.legend.Legend at 0x7ff48b3c3350>
```



```
In [ ]: test_loss, test_acc = model2.eval(loader=test_loader, BATCH_SIZE=100)
print("Test loss for Model 2 is ", test_loss)
print("Test accuracy for Model 2 is ", test_acc)
```

```
Test loss for Model 2 is  0.010438133199435765
Test accuracy for Model 2 is  0.86
```

Model 2 - Performance Analysis

Looking at the training and validation set accuracy, we can see that there is no overfitting happening. The test accuracy for this model is 85% which is slightly better than model 1. One constraint that we considered in network design is the complexity of the model since going for model of higher complexity would result in training time increasing a lot which becomes an overhead in performing iterations. Also, complex models are more prone to overfit the data.